

Figure 4: Overall Top-K Accuracy Comparison (210 runs, 7 anomaly types $\times$ 3 pipeline sizes)

